

A reusable RAG pipeline


A BioASQ reproduction kit

A model-organism case study

Yun Wang, Blaž Zupan · University of Ljubljana, Slovenia



This poster is available digitally. Scan me.



Top-ranked retrieval in BioASQ 14b Phase A

Ranked 1st in document retrieval in 3 of 4 batches · preliminary results

01 A reusable RAG pipeline github.com/fulaibaowang/RAG-scripts

Configurable

Pipeline stages and settings in one config

Reproducible


Docker image, pinned dependencies

Portable

BioASQ, TREC-RAG, your corpus

Agent-friendly

Built-in guidance for AI coding assistants



02 A BioASQ reproduction kit

github.com/fulaibaowang/BioASQ


Scripts and configurations to reproduce our BioASQ retrieval runs.



One batch of 80 questions · one A100 GPU

16 CPU cores **~80 GB** RAM **~1.5 h** retrieve + rerank **~150 GB** one-time: corpus, indexes, container

Automatic diagnostic plots · detailed analysis in the working note



Examples on BioASQ 13b, batches 1–4 averaged

Email me
yun.wang@fri.uni-lj.si



or connect
on LinkedIn

FRI
UNIVERSITY OF LJUBLJANA
Faculty of Computer and Information Science

SMASH
machine learning for science and humanities postdoctoral program

BEYOND BIOASQ

03 A model-organism case study

github.com/fulaibaowang/dictycite

Why is our system called DictyCite?

Dicty (*Dictyostelium*) is a model organism in biology. We built a literature retrieval benchmark and case study for it from curated biological annotations.

Benchmarking Literature Retrieval for a Model Organism: A Dictyostelium Case Study

Our *Dictyostelium* case study suggests that tailoring literature retrieval to a biological domain can be worthwhile, using curated annotations to enrich queries and choosing retrievers and rerankers that fit the task.

Accepted at Discovery Science 2026

Mainz, Germany · 5–9 October 2026